

Build dependency graph — what parallelizes, what cannot, and why

Written 2026-08-07. Companion to `docs/IMPLEMENTATION-PLAN.md`, which says *what* to build and how long each step takes. This document says **what may run at the same time**, and it is written to be handed to an orchestrating agent.

The answer up front: the critical path is **13.3 hours** with correct parallelization, against **~36 h** if the work is run as written today and **17.3 h** if it is parallelized naively. The irreducible floor is **8.41 h** of pure job time — so only **~4.9 h** of code and setup can ever sit ahead of the jobs. **Everything beyond that is either waiting on a human or wasted.**


1. The three currencies, which are not interchangeable

Most parallelization mistakes here come from mixing these up.

currency	unit	what buys more of it	what caps it
agent-hours	code items	more agents in more worktrees	file contention (§3), not logic
wall-hours	AWS jobs	more array children, more workers	128 vCPU compute environment
calendar	human gates	nothing — you wait	approval latency, license acceptance

A code item and a job of the same nominal duration are **not** substitutable: 8 agent-hours of code can run alongside 8 wall-hours of jobs, but two code items touching the same function cannot run alongside each other at all.

2. The graph



Red = critical path. Blue = parallel. Grey = deferrable to a second image. Amber = waits on a human.

3. ⚠️ The real constraint is file contention, not logic

This is the part an orchestrator gets wrong. Eleven code items look independent on a task list. They are not — five of them edit `corpus_build.py`, and within that file they cluster into **two functions**:

function	items that edit it	verdict
<code>_reader_for</code>	C1, C3, C11	one owner, or three-way conflict
<code>run_bundle</code>	A2b, C7	one owner
<code>corpus.py</code>	B6, C3	different constants, but same file — merge risk

file	items	contention
<code>corpus_build.py</code>	C1, C3, C4→A2b, C7, C11	5 — hottest file in the plan
<code>corpus.py</code>	C3, B6	2
everything else	1 each	none

Rules for the orchestrator:

1. **One agent owns one function, never one file.** `_reader_for` and `run_bundle` can be two agents in two worktrees, because they do not overlap textually — but both must be told the other exists.
2. **Distinct-file items are genuinely free.** B1, B2, B3, B4, B5 can be five simultaneous agents with zero coordination.
3. **Worktrees are mandatory** for anything touching `corpus_build.py`, per this repo's own convention:
`../Capstone_LLM-worktrees/edullm-data/<agent-id>--<task-slug>` on `agent/<agent-id>/<task-slug>`.
4. **Merge order is** `_reader_for` → `run_bundle` → **everything else**, because the first two are on the critical path and a late merge conflict there costs image-build time.

4. The four hard serialization points

Nothing parallelizes past these. They are why 13.3 h is the floor and not 8.41 h.

S1 — the image build gates every AWS job

Container images build **only from** `edullm/**` branches (`.github/workflows/edullm-platform-build.yml:9-10`). A merge to `main` builds **nothing, silently**, and a submission naming that commit is refused for having no image. So **all code that the jobs need must land in one push**, and every job waits on it.

Consequence: batching code into one image is *better* than shipping two. Measured against the graph, a two-image scheme (pre-pass early, build later) came out **0.1 h worse** — the second build's 0.5 h lands directly on the critical path.

S2 — `FREEZE` is a human decision with an unbounded duration

Adding a source after the plan is generated renames **98% of shards** and voids **882B tokens** (`IMPLEMENTATION-PLAN.md §0`). So the mix must be final first. **Its duration is a decision, not a computation** — it is the one node whose length no amount of parallelism touches.

S3 — `BUILD` is capped at 128 vCPU

6.6 h is the tokenize floor at the cap. Slicing bundles does not lower it; it only prevents a long tail (the 159B-token `stackv2-edu` bundle is a 4.37 h single child unless file-sharded).

S4 — `VD1` cannot start until `PUB1` finishes

`verify --deep` re-hashes published objects, so it is strictly after publish. At 8 workers it is 1.49 h, and it sits on the critical path with no way around it. At `--hash-workers 1` it is **11.7 h and exceeds its timeout**, which is the single highest-value flag in the plan.

5. The critical path, and everything with slack

Critical path — 13.31 h:

from → to	node	why it cannot move
0.00 → 4.00	A2a hash pre-pass driver	nothing to parallelize with; <code>PASS1</code> needs it
4.00 → 4.50	IMG image build	S1
4.50 → 4.90	SMOKE live-HF smoke test	mandatory; the path has never run
4.90 → 11.50	BUILD ~100 bundles	S3, the 128 vCPU floor
11.50 → 11.80	PUB1 publish stage 1	after build
11.80 → 13.29	VD1 verify --deep stage 1	S4
13.29 → 13.31	PR1 promote stage 1	after both gates

Everything else has slack and should be started at t=0 regardless:

node	duration	slack	note
M1 bandwidth	0.2 h	large	but run it first anyway — it calibrates every other estimate
M2 / M4 samples	1.0 h each	~3 h	gate <code>FREEZE</code> , not the build
M3 footer count	0.1 h	~3.9 h	blocked on a human licence acceptance (G1)
A2b keep-list consumer	4 h	0 h — joint-critical	must finish before IMG, same as A2a
B3 thread Gate A	4 h	7.5 h	only <code>GA1</code> / <code>GA2</code> need it
B5 rebuild decon index	4 h	~0.9 h — nearly critical	gates <code>FREEZE</code> ; start it early
B1 / B2 / B4	≤1 h	large	trivially parallel
C3 / C11	5 h combined	∞ — deferrable	pure optimization + instrumentation. Do not put them in the first image

Stage 2 is entirely parallel to stage 1's tail. `PUB2` / `GA2` / `VD2` / `PR2` total 0.25 h and finish while stage 1 is still verifying, so stage 2 contributes **nothing** to the critical path.

6. What the orchestrator should actually launch

Wave 0 — immediately, 8 concurrent workstreams

stream	work	kind
1	M1 bandwidth measurement	job — do this first, it recalibrates the rest
2	M2 Dolma3 sample + M4 doc lengths	job
3	Ask the owner to accept the Nemotron licence gate	human
4	A2a hash pre-pass driver (owns <code>corpus_filter.py</code>)	agent, worktree
5	A2b keep-list consumer (owns <code>run_bundle</code>)	agent, worktree
6	C1 + B6 (owns <code>_reader_for</code> + <code>SHARD_TOKENS</code>)	agent, worktree
7	B5 rebuild the decon index (external repo)	agent
8	B1 + B2 + B3 + B4 (four distinct files)	agent(s)

Do not launch C3 or C11 in wave 0. They are deferrable, they contend with C1 on `_reader_for` , and including them adds ~5 h of agent time to the critical path for zero wall-clock benefit.

Wave 1 — merge, then one push

Merge in order `_reader_for` → `run_bundle` → the rest. Push to an `edullm/**` branch. **One image.**

Wave 2 — the gate

`FREEZE` the mix, generate the plan, stage the sources. `STAGE` can overlap `SMOKE` .

Wave 3 — jobs

PASS1 → BUILD (16 children × 8 vCPU) → then stage 1 and stage 2 publish/validate **in parallel**.

Every AWS job needs a separate human release. Nothing auto-publishes.

7. Where parallelism is wasted — spend nothing here

tempting	why it does not help
More than 16 build children	128 vCPU ÷ 8 = 16. More children just queue
Splitting BUILD further	6.6 h is the CPU floor at the cap, not a scheduling artifact
Two container images	measured 0.1 h worse — the second build lands on the critical path
Parallelizing STAGE harder	it has ~4 h of slack; it is never the constraint
Doing C3 / C11 now	deferrable, and they contend on the hottest function in the repo
More agents on corpus_build.py	file contention. A third agent there produces conflicts, not speed
Rushing before M1 lands	every read estimate is calibrated on an unmeasured bandwidth

8. The orchestrator brief — copy this

Hand this to the agent that runs wave 0. It encodes the constraints above as rules rather than prose.

You are orchestrating the Phase-0 build for `pretrain/final-dataset` . Launch these 8 streams concurrently and hold them to these rules.

Stream 1 (job, FIRST): measure in-region S3 read bandwidth and HF CDN throughput. Report both. Every read estimate in `IMPLEMENTATION-PLAN.md §8A` is calibrated on an unmeasured ~ 85 MB/s borrowed from an S3 measurement; one plausible reconciliation says the CDN is ~ 8.4 MB/s. If it is, tell me before anything else proceeds — it changes the staging decision.

Stream 2 (job): Dolma3 adult-content prevalence at **random offsets** (a prior attempt could not separate signal from HuggingFace preview ordering), plus mean document length for the 5 unmeasured stage-2 sources.

Stream 3 (human): ask the owner to accept the Nemotron-CC-Math licence gate. Its text and id columns are **UNVERIFIED** and it is the only stage-1 source whose text column we cannot name.

Streams 4–6 (code, one worktree each, one function each — NOT one file each): - **4:** the `flat_np.uint64` hash pre-pass. Owns `corpus_filter.py` . - **5:** the keep-list consumer. Owns `run_bundle` in `corpus_build.py` . - **6:** the FinePhrase id partition + the shard-size constant. Owns `_reader_for` in `corpus_build.py` and `SHARD_TOKENS` in `corpus.py` .

⚠ **Streams 5 and 6 both edit `corpus_build.py` in different functions.** Tell each about the other. Worktree convention: `../Capstone_LLM-worktrees/edullm-data/<agent-id>--<task-slug>` on `agent/<agent-id>/<task-slug>` . Merge order is **stream 6** → **stream 5** → **the rest**.

Stream 7 (code): rebuild the decontamination index from **raw** benchmark fields (question alone, question + each choice, question + correct answer) **in addition to** the rendered form. External repo, no conflict. **It gates the mix freeze — do not let it start late.**

Stream 8 (code, four distinct files, no coordination needed): pin tokenizers in `pyproject.toml` ; fix the boundary-marker prefix guard in `corpus_pack.py` **and** the test that asserts the table length is 1; thread the Gate A profile checks in `pretrain_tokens_v1.py` **and** raise `max_pool_connections` in `s3.py` (threading against a 10-connection pool self-throttles); drop `data_provenance_initiative` from the registry.

Do NOT start: `val-bundle` file-sharding or the `bytes_fetched` wiring. Both are deferrable, both contend with stream 6 on `_reader_for` , and including them adds ~ 5 agent-hours to the critical path for **zero** wall-clock gain.

When all streams land: merge in the stated order, push **once** to an `edullm/**` branch — images build only from that namespace, a merge to `main` builds nothing silently — and stop. **Do not submit any AWS job.** Report what landed and wait for a release decision.

Standing rules: write findings to disk incrementally (a prior wave lost 4 agents to an API budget cap and only on-disk partials survived); grade every number MEASURED / DERIVED / CARD / UNVERIFIED; "*nobody has measured this*" is a finding, not a failure; and if you find that one of my numbers is wrong, say so — two of them already were.

9. Honest limits of this graph

- **Durations for code items are estimates, not measurements.** The *ordering* is well-evidenced; the absolute agent-hours are my judgement. An earlier session's per-bundle ETAs were all retracted by their own author with the verdict "*every per-bundle ETA I gave was wrong, in both directions.*"

- BUILD 's **6.6 h** assumes **linear vCPU scaling**, measured only at 32 vCPU. Never tested at 128.
- Every read duration assumes **~85 MB/s borrowed from an S3 measurement**. If the HF CDN is really ~8.4 MB/s, STAGE moves from 0.5 h to days and becomes critical. **M1 exists to settle exactly this — which is why it is wave 0, stream 1.**
- The **13-gram decontamination scan's CPU cost has never been measured** (~193 billion Python-level blake2b calls). If it is comparable to tokenization, BUILD is longer than 6.6 h.
- FREEZE and G1 have no duration here because they are human. In practice they may dominate the calendar, and no amount of parallelism touches them.