

Task register — final-dataset

Written 2026-08-07. The #NN ids are referenced throughout IMPLEMENTATION-PLAN.md , BUILD-DEPENDENCY-GRAPH.md and HANDOFF-FINAL-DATASET.md . They previously lived **only in a session task tool**, so a fresh agent could read #21 in three documents and find no definition anywhere. This file is the definition. **It is the crosswalk between all three numbering schemes** — task id, graph node, and Phase 0 item.

Nothing here has been submitted to AWS. Every job needs a platform submission and a human release.

Wave 0 — the critical-path code

#	task	graph	Phase 0	changes the plan?	est
#28	File-shard the BIG bundles. --shard/--of strides <i>bundles</i> (corpus_build.py:676), so DCLM's 410B lands in one child = 10.85 h even on a whole 32-vCPU instance . Needs plan-time ordinal ranges. The longest item and the new critical path	C3b	3b	YES	1–2 d
#22	Replace the dedup set with a flat np.uint64 pre-pass. The set needs 27.92 GB for DCLM in a 15.03 GB container (§5.2a)	A2a + A2b	4	no	1 d
#21	Wire the FinePhrase id partition into _reader_for . Cannot be retrofitted after tokenization. Ships with the budget correction or not at all	C1	1	YES	~5 lines
#9	Decide SHARD_TOKENS . Code says 25,001,984; the report now agrees. Confirm and stop carrying two values	B6	11	YES	1 h
#24	Rebuild the decon index from raw benchmark fields, not 5-shot renders. Gates FREEZE	B5	12	no	4 h

Wave 0 — parallel, distinct files, no coordination

#	task	graph	Phase 0	est
#23	Pin tokenizers in pyproject.toml . It is imported at corpus_build.py:631 and declared nowhere	B1	5	1 line
#10	Thread the Gate A profile checks + raise max_pool_connections . ~100 lines, not the ~20 an earlier draft said	B3	6	1 d
#16	Drop data_provenance_initiative — ships GSM8K in Flan CoT format at 6 repeats. 0.51% of tokens	B4	9	registry edit
—	Fix the boundary-marker prefix guard in corpus_pack.py , and the test asserting the table length is 1	B2	8	~5 lines
—	Record FilterStats in the receipt. Until then no dedup claim is auditable	—	7	~10 lines
#11	Query the live validator timeout. edullm-validator:12 's is recorded nowhere. Read-only	—	10	1 call

Wave 0 — measurements

#	task	graph	est
#26	Measure in-region S3 + HF CDN bandwidth. RUN THIS FIRST — every read estimate borrows ~85 MB/s from an S3 measurement, and one reconciliation implies the CDN is ~8.4 MB/s	M1	10 min
#14	Dolma 3 adult-content prevalence, sampled at random offsets — a prior attempt could not separate signal from HuggingFace preview ordering. Blocking for that source	M2	1 h
—	~~Nemotron-CC-Math's real dolma2 count~~  DONE 2026-08-07 — 134.0B (472,213,218,716 bytes × 0.283686 tok/byte, 1,920 random-offset docs, seed 42; 3 ≈ 83.6B + 4plus ≈ 50.4B). Measured by a teammate with gate access. This closed the M3 node, NOT task #17	M3	done
#17	Re-run Marin's 13-gram contamination scan on Nemotron-CC-Math. A different task from the count above, and still open. The original scan found 11,868 contaminated documents against MATH500 alone — 13.2× the publisher's entire reported removal budget — including verbatim GSM8K test items at Jaccard 1.0. Now unblocked: the gate is accepted and 107+ GiB is already staged in S3, so this runs in-region against staged bytes	—	1 job
—	Mean doc length for 5 unmeasured stage-2 sources. The dolma3 QA source is the one plausibly near the 20-token EOS floor	M4	1 h

Deferrable — do NOT put in the first image

#	task	graph	why deferred
#25	File-shard the VAL bundles. 43% of bytes moved serves 0.39% of tokens. NOT the same item as #28 — this one is a read-volume saving, #28 is a wall-clock prerequisite	C3	∞ slack; contends with #21 on _reader_for
—	Wire bytes_fetched through from corpus_read to the build	C11	instrumentation only

Then, in order

#	task	note
#20	Freeze the full two-stage plan. THE REAL GATE. Its duration is a decision, not a computation	after every "changes the plan" item lands
#6	Ingest. Stage → smoke test → build in waves → publish two datasets → Gate A → verify --deep → promote	
#1	Register reservoir-dolma2 with the platform	independent of this corpus

Downstream of the corpus — training-side, not build-side

#	task
#19	Fix <code>MoELoadBalancingLossGranularity</code> before any training run. Cannot be annealed in — switching at 10% recovers ~55%
#7	Run probe-fitted mixture selection before the real run
#18	Use BPB, not accuracy, for any ablation at this scale — MCF is random below ~400B tokens
#8	Store MTLD as a per-document label at ingest — labels are inside <code>manifest_sha256</code> , so unbackfillable
#13	Settle SuperBPE empirically with the existing byte-matched A/B corpora (~\$600)
#12	Fix the missing EOS in the SuperBPE tokenizer. Affects two already-published corpora, NOT this one — dolma2's <code>eos_token_id</code> is 100257 and its check runs

Closed

#	task	outcome
#5	Measure the real FinePhrase overlap	0.2683 distinct on a complete-column read of 287,000 ids
#15	Interleave sources within shards	Not expressible in data. <code>labels</code> is one dict per shard path and Gate A compares by full dict equality. Trainer-side (#19)
#4	Wire <code>keeps_id</code> into the build path	superseded by #21, which is the same fix scoped to the reader
#27	Fix 23 verified doc inconsistencies	this pass